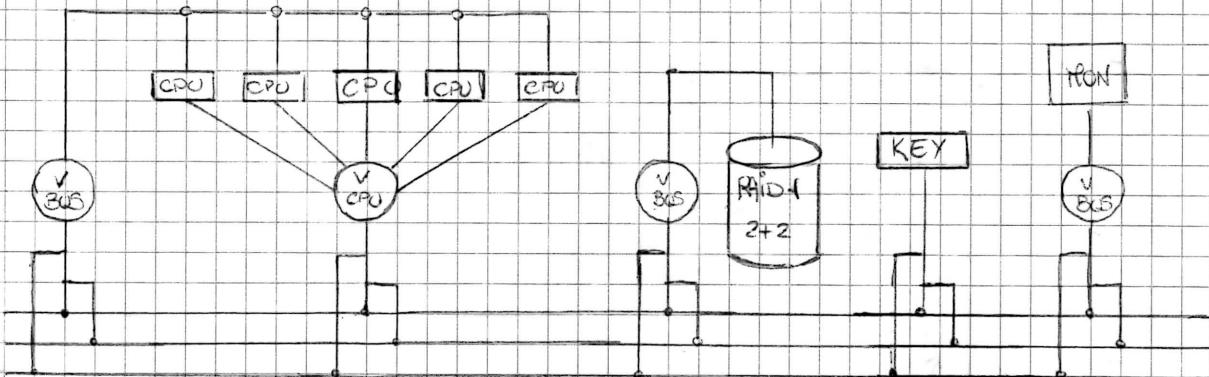
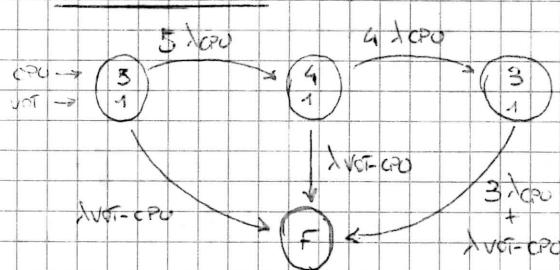


Exercise 1 Hypothesis: 3 switch off the system when a component faults.



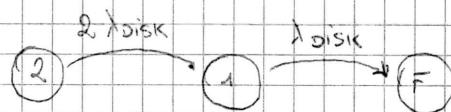
RELIABILITY

- CPU + VOTER CPU



$$R_{CPU}(t) = 1 - P_F(t)$$

- RAID (single pair)



$$R_{pair} = 1 - P_F(t) \Rightarrow R_{RAID} = (R_{pair})^2$$

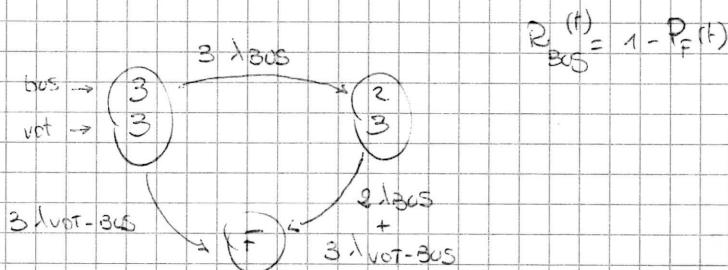
$$\left\{ \begin{array}{l} P_{5,1}(t) = -P_{5,1}(t)[5\lambda_{CPU} + \lambda_{VOT-CPU}] \\ P_{4,1}(t) = +P_{5,1}(t)[5\lambda_{CPU}] - P_{6,1}(t)[4\lambda_{CPU} + \lambda_{VOT-CPU}] \\ P_{3,1}(t) = +P_{6,1}(t)[4\lambda_{CPU}] - P_{7,1}(t)[3\lambda_{CPU} + \lambda_{VOT-CPU}] \\ P_F(t) = +\lambda_{VOT-CPU} [P_{5,1}(t) + P_{6,1}(t) + P_{7,1}(t)(3\lambda_{CPU})] \\ P_{5,1}(t) + P_{6,1}(t) + P_{7,1}(t) + P_F(t) = 1 \\ P_{5,1}(0) = 1 \end{array} \right.$$

$$R_{RAID} = 1 - (1 - R_{disk})^2$$

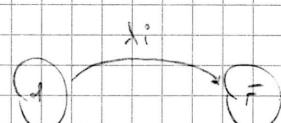
at least one disk has to work
in the pair (no one of the two disks
fails)

$$R_{RAID} = [R_{pair}]^2 \quad ? \quad ? \text{ has to be valid for both of the pair}$$

- BUS + VOTER BUS



- KEYBOARD / MONITOR



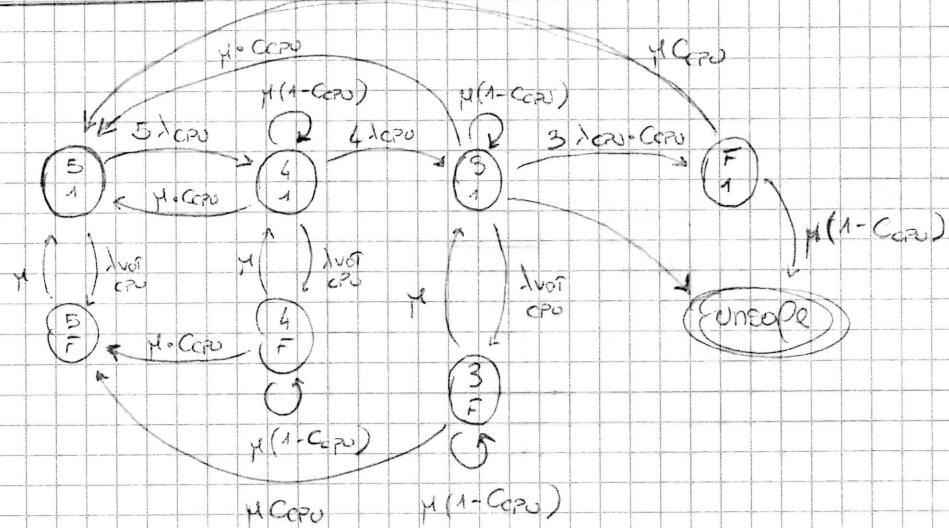
$$R_{i,i}(t) = 1 - P_F(t) \quad \text{where } i = \text{key / mon}$$

Finally:

$$R_{tot}(t) = R_{CPU}(t) \cdot R_{RAID}(t) \cdot R_{BUS}(t) \cdot R_{key}(t) \cdot R_{mon}(t)$$

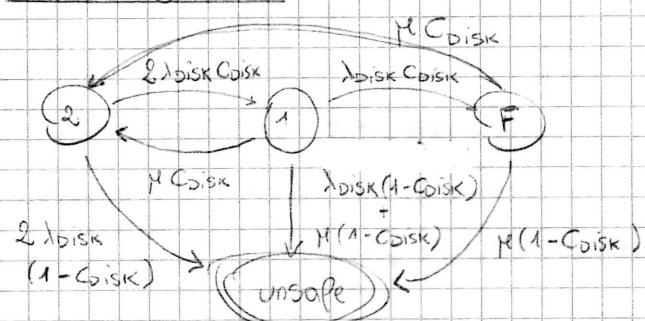
Safety Hypothesis: The coverage factor of the voters is 1.

- CPU + VOTER CPU



$$S_{CPU}(t) = 1 - P_{unsafe}(t)$$

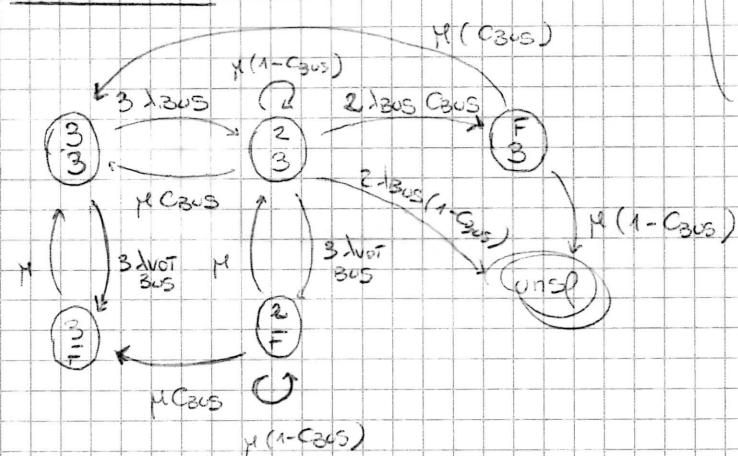
- RAID (single pair)



$$\begin{aligned} P_2(t) &= -P_2(t) [2 \lambda_{disk} [C_{disk} + (1 - C_{disk})]] + \\ &\quad \mu_{disk} (P_F(t) + P_1(t)). \\ P_1(t) &= -P_1(t) \{ \lambda_{disk} [C_{disk} + (1 - C_{disk})] \} + \\ &\quad \mu [C_{disk} + (1 - C_{disk})] + P_2(t) (2 \lambda_{disk} C_{disk}) \\ P_F(t) &= -P_F(t) [\mu (C_{disk} + (1 - C_{disk}))] + P_1(t) (\lambda_{disk} C_{disk}) \\ P_{unsafe}(t) &= P_2(t) [2 \lambda_{disk} (1 - C_{disk})] + P_1(t) [\lambda_{disk} (1 - C_{disk}) + \\ &\quad \mu (1 - C_{disk})] + P_F(t) [\mu (1 - C_{disk})] \\ P_2(t) + P_1(t) + P_F(t) + P_{unsafe}(t) &= 1, \quad P_2(0) = 1 \end{aligned}$$

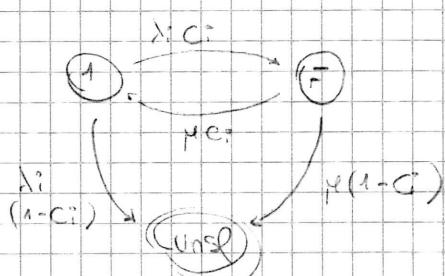
$$S_{PAIR}(t) = 1 - P_{unsafe}(t) \Rightarrow S_{RAID}(t) = (S_{PAIR})^2$$

- BUS + VOTER BUS



$$S_{BUS}(t) = 1 - P_{unsafe}(t)$$

- KEYBOARD / MONITOR



$$S_i(t) = 1 - P_{unsafe}(t) \quad \text{where } i = \text{Key / mon}$$

Finally:

$$S_{tot}(t) = S_{CPU}(t) \cdot S_{RAID}(t) \cdot S_{BUS}(t) \cdot S_{KEY}(t) \cdot S_{MON}(t)$$

Exercise 2

20.02.2014

L'NVA (Network Value Analysis) è un metodo che ci permette di calcolare gli indici di performance (response time medio, throughput, lunghezza della coda, ecc...) per reti chiuse.

Per "rete chiusa" intendiamo una rete che è formata da un numero limitato e fisso di utenti quindi ha un limite sul n° di richieste che il sistema deve soddisfare.

(È composta da due fasi: richiesta del servizio e thinking time).

L'NVA sfrutta un metodo iterativo basato su tre equazioni e su un teorema.

Definizioni:

X_i = throughput medio della rete di coda

V_i = n° medio di visite per richiesta, alla coda i

S_i = service time medio per richiesta, sul server i

$W_i^{(n)}$ = waiting time, alla coda i

= $n_i^{(n)} \cdot S_i$, dove $n_i^{(n)} = n$ medio di richieste fatte nella coda i da una richiesta i

$R_i^{(n)}$ = residence time medio per una richiesta, alla coda i

$$= S_i + W_i^{(n)} = S_i + [n_i^{(n)} \cdot S_i] = S_i (1 + n_i^{(n)})$$

$R_i^{(n)}$ = residence time totale per una richiesta alla coda i considerando tutte le visite a tale coda

$$= V_i \cdot R_i^{(n)}$$

D_i = media del service time totale per una richiesta alla coda i considerando tutte le visite alla coda

$$= V_i \cdot S_i$$

$R(n)$ = response time medio della rete di coda = $\sum D_i^{(n)}$

Teorema:

Il n° medio di richieste che una nuova richiesta entrante vede nella coda i-esima nella rete di coda quando già ci sono n richieste (cioè $n_i^{(n)}$), è uguale al n° medio di richieste che ha nella coda i-esima quando ce ne erano n-1 di richieste (cioè $n_i^{(n-1)}$).
In rete di coda.

$$\text{Quindi: } n_i^{(n)} = n_i^{(n-1)}$$

Allora, da questo teorema, otteniamo:

$$R_i^{(n)} = S_i [1 + n_i^{(n-1)}]$$

Riplicando entrambi i membri per V_i , otteniamo la 1^a equazione dell'RNA.

$$R_i^1(n) = D_i [1 + n_i(n-1)] \quad \text{per risorse in coda} \quad \left. \begin{array}{l} \\ \end{array} \right\} \text{Residence time equation}$$

or $= D_i$ per risorse di ritardo

Applicando la legge di Little alla intera rete di code ($n = X_0(\cap T R_i(n))$) otteniamo la 2^a:

$$X_0(n) = \frac{n}{R_0(n)} \quad \text{Throughput equation where } R_0(n) = \sum_{i=1}^n R_i^1(n) \quad \text{d} = n^{\circ} \text{ of the queue}$$

La 3^a equazione dell'RNA è ottenuta combinando la legge di Little e quella del processo forzato:

$$n_i(n) = X_0(n) \cdot R_i^1(n)$$

Ma come faccio a calcolarmi i valori di $R_i^1(n)$ se questo dipende da $n_i(n-1)$?

A sua volta $n_i(n-1)$ dipende da $R_i^1(n-1)$ che dipende da $n_i(n-2)$, ecc...

Dobbiamo partire da $n=0$ e calcolare fino al valore di n a cui siamo interessati (i.e. n° di richieste massima che possono essere soddisfatte dalla rete).

Per fortuna abbiamo che:

$$n_i(0) = 0 \quad \forall i$$

Questo perché quando non ci sono richieste nel sistema, allora la lunghezza delle code è nulla per tutti i componenti!

Quindi, a questo punto cominciamo a calcolare.

$$n_i(0) \rightarrow R_i^1(1) \rightarrow X_0(n) \rightarrow n_i(1) \rightarrow R_i^1(2) \rightarrow X_0(2) \rightarrow n_i(2) \rightarrow \dots$$

Calcolati tutti i valori, potrei potremmo utilizzarli per calcolare:

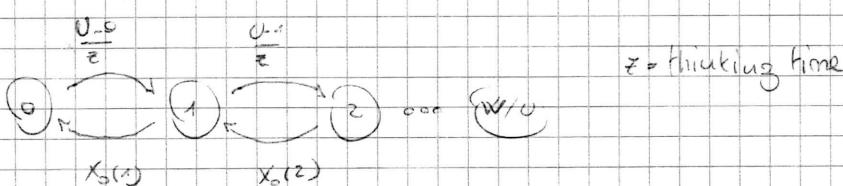
$$\bar{x} = \sum_{i=1}^{W/0} p_i \cdot X_0(i) \quad \text{Average throughput of the networkic}$$

$\bar{N} = \sum_{i=0}^W p_i \cdot i \quad \text{average n° of users in the queue networkic}$

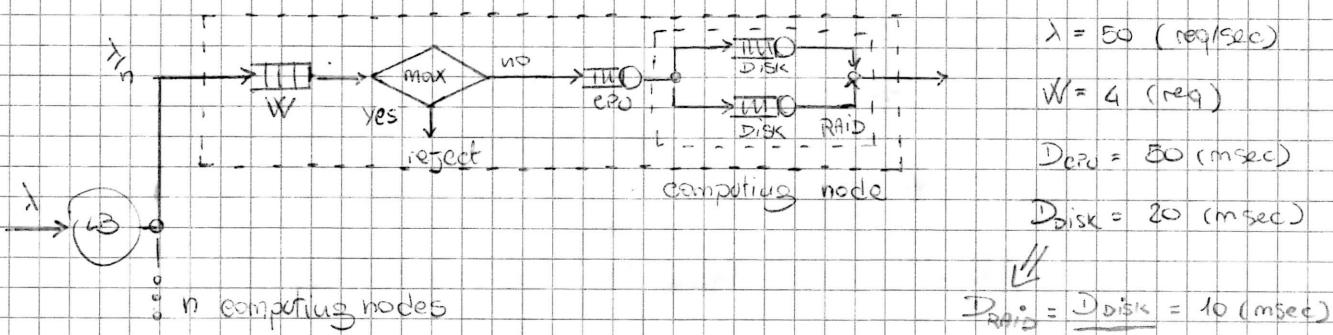
$\bar{R} = \frac{\bar{N}}{\bar{x}} \quad \text{average service time}$

$W = \text{Lunghezza coda degli utenti (se c'è)}$
oppure
 $U = n^{\circ} \text{ utenti del sistema}$

dove p_i rappresenta la probabilità di trovarsi nello stato in cui nel sistema ci sono i -richieste.



$\tau = \text{thinking time}$

Exercise 3

3 hypothesis: 3 hypothesize that there is a load balancer which equally distributes the load of the request to the n servers.

3 have to calculate the n° of computing nodes that allows me to have the fraction of lost requests less than 10%.

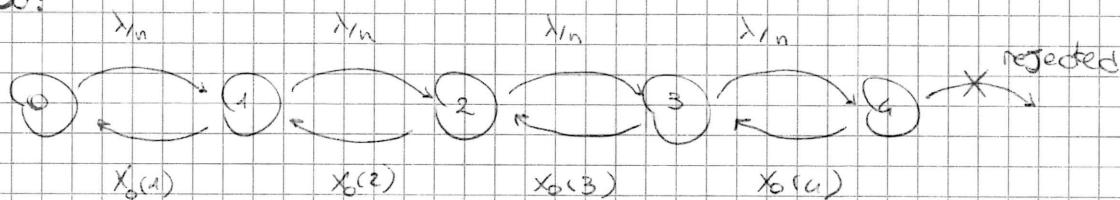
There aren't the R_{CPU} , R_{Disk} parameters so 3 don't have to consider the cases of faulty servers.

SINGLE SERVER SCENARIO

The arrival rate of the requests, due to the load balancer, is $\frac{\lambda}{n}$ and it's independent of the n° of requests already inside the server.

Instead, the throughput of the server is dependent on that.

So:



For the ~~flow-in = flow-out principle~~:

$$\left\{ \begin{array}{l} P_0 \cdot \frac{\lambda}{n} = P_0 \cdot X_0(1) \\ P_1 \cdot \frac{\lambda}{n} = P_2 \cdot X_0(2) \\ P_2 \cdot \frac{\lambda}{n} = P_3 \cdot X_0(3) \\ P_3 \cdot \frac{\lambda}{n} = P_4 \cdot X_0(4) \\ \sum_{i=0}^4 P_i = 1 \end{array} \right. \quad \begin{array}{l} P_0 = P_0 \left[\frac{\lambda}{n} \cdot \frac{1}{X_0(1)} \right] \\ P_2 = P_1 \left[\frac{\lambda}{n} \cdot \frac{1}{X_0(2)} \right] = P_0 \left[\left(\frac{\lambda}{n} \right)^2 \cdot \frac{1}{X_0(1)} \cdot \frac{1}{X_0(2)} \right] \\ P_3 = \dots = P_0 \left[\left(\frac{\lambda}{n} \right)^3 \cdot \frac{1}{X_0(1)} \cdot \frac{1}{X_0(2)} \cdot \frac{1}{X_0(3)} \right] \\ P_4 = \dots = P_0 \left[\left(\frac{\lambda}{n} \right)^4 \cdot \frac{1}{X_0(1)} \cdot \frac{1}{X_0(2)} \cdot \frac{1}{X_0(3)} \cdot \frac{1}{X_0(4)} \right] \end{array}$$

To calculate the throughput of the server, (based on the n° of request inside ~~if~~ it), 3 have to use the EVA algorithm.

The iterative procedure goes from n=1 to n=4.

We know that: $n_i(\emptyset) = 0 \forall i$ {where $i = CPU, disk$ }

So, using the 3 equations:

$$R_i(n) = D_i [1 + n_i(n-1)], \quad X_0^{(n)} = n / \sum R_i(n), \quad n_i(n) = X_0^{(n)} \cdot R_i(n)$$

$$R_{CPU}^1(1) = D_{CPU} [1 + n_{CPU}(1)] = 50 \text{ (msec)}$$

$$R_{RAID}^1(1) = D_{RAID} [1 + n_{disk}(1)] = 10 \text{ (msec)} \rightsquigarrow \text{Having a RAID 1 or 3 use a disk or 2 use the other one.}$$

$$R_0^1(1) = R_{CPU}^1(1) + R_{RAID}^1(1) = 60 \text{ (msec)}$$

$$X_0(1) = 1/R_0^1(1) = 1/(60 \cdot 10^{-3}) \text{ (sec)} = 16,67 \text{ (req/sec)}$$

$$n_{CPU}(1) = X_0(1) \cdot R_{CPU}^1(1) = 16,67 \cdot 50 \text{ (req)} = 0,83 \text{ (req)} \quad \left. \begin{array}{l} \\ \end{array} \right\} \begin{array}{l} n_{CPU}(1) + D_{RAID}(1) = 1 \text{ OK!} \\ \end{array}$$

$$n_{disk}^1(1) = X_0(1) \cdot R_{RAID}^1(1) = 16,67 \cdot 10 = 0,17 \text{ (req)}$$

(n=2)

$$R_{CPU}^1(2) = D_{CPU} [1 + n_{CPU}(1)] = 50 \text{ (msec)} [1,00] = 50,00 \text{ (msec)} = 0,0000 \text{ (sec)}$$

$$R_{RAID}^1(2) = D_{RAID} [1 + n_{disk}^1(1)] = 10 [1,00] = 10,00 \text{ (msec)} = 0,0000 \text{ (sec)}$$

$$R_0^1(2) = R_{CPU}^1(2) + R_{RAID}^1(2) = 50,00 + 10,00 = 60,00 \text{ (msec)}$$

$$X_0(2) = 1/R_0^1(2) = 1/60,00 = 17,38 \text{ (req/sec)}$$

$$n_{CPU}(2) = X_0(2) \cdot R_{CPU}^1(2) = 17,38 \cdot 50,00 = 0,86 \text{ (req)} \quad \left. \begin{array}{l} \\ \end{array} \right\} \Sigma = 2 \text{ OK}$$

$$n_{disk}^1(2) = X_0(2) \cdot R_{RAID}^1(2) = 17,38 \cdot 10,00 = 0,46 \text{ (req)}$$

(n=3)

$$R_{CPU}^1(3) = D_{CPU} [1 + n_{CPU}(2)] = 50 \text{ (msec)} [2,00] = 50,00 \text{ (ms)} = 0,125 \text{ (sec)}$$

$$R_{RAID}^1(3) = D_{RAID} [1 + n_{disk}^1(2)] = 10 [2,00] = 20,00 \text{ (ms)} = 0,0000 \text{ (sec)}$$

$$R_0^1(3) = R_{CPU}^1(3) + R_{RAID}^1(3) = 50,00 + 20,00 = 70,00 \text{ (msec)}$$

$$X_0(3) = 1/R_0^1(3) = 1/70,00 = 14,29 \text{ (req/sec)}$$

$$n_{CPU}(3) = X_0(3) \cdot R_{CPU}^1(3) = 14,29 \cdot 50,00 = 0,71 \text{ (req)} \quad \left. \begin{array}{l} \\ \end{array} \right\} \Sigma = 3 \text{ OK!}$$

$$n_{disk}^1(3) = X_0(3) \cdot R_{RAID}^1(3) = 14,29 \cdot 10,00 = 0,46 \text{ (req)}$$

(n=4)

$$R_{CPU}^1(4) = D_{CPU} [1 + n_{CPU}(3)] = 50 \text{ (msec)} [3,00] = 50,00 \text{ (msec)} = 0,125 \text{ (sec)}$$

$$R_{RAID}^1(4) = D_{RAID} [1 + n_{disk}^1(3)] = 10 [3,00] = 30,00 \text{ (msec)} = 0,0000 \text{ (sec)}$$

$$R_0^1(4) = R_{CPU}^1(4) + R_{RAID}^1(4) = 50,00 + 30,00 = 80,00 \text{ (msec)}$$

$$X_0(4) = 1/R_0^1(4) = 1/80,00 = 12,50 \text{ (req/sec)}$$

Now, 3 can calculate:

$$X(0) = \sum_{i=1}^4 p_i \cdot X_0(i) = [p_1 \cdot X_0(1)] + [p_2 \cdot X_0(2)] + [p_3 \cdot X_0(3)] + [p_4 \cdot X_0(4)]$$

↑ average throughput of a single server/computing node

$$N(0) = \sum_{i=0}^4 p_i \cdot i = 0 + p_1 + 2p_2 + 3p_3 + 4p_4$$

↑ average n° of requests into the single server

$$R(0) = \frac{N(0)}{X(0)}$$

(from the Little's Law)

↑ response time of the single server.

And, for the whole system (composed of n servers):

$$\bar{X} = n \cdot X(0) \leftarrow \text{web site throughput}$$

$$\bar{R} = R(0) \leftarrow \text{web site response time}$$

Finally, the fraction of rejected requests is given by:

$$\frac{\lambda - \bar{X}}{\lambda} < 10\% \quad \leftarrow \text{This concerns the whole web server!}$$

If instead, the professor meant about the single computing node, then the fraction of lost requests < 10% is given by:

$$p_4 < 0,01$$

(Where p_4 is the probability to stay in the case in which the computing node has 4 request in queue, that is the maximum supported.)

Beyond 4 the request is rejected.)